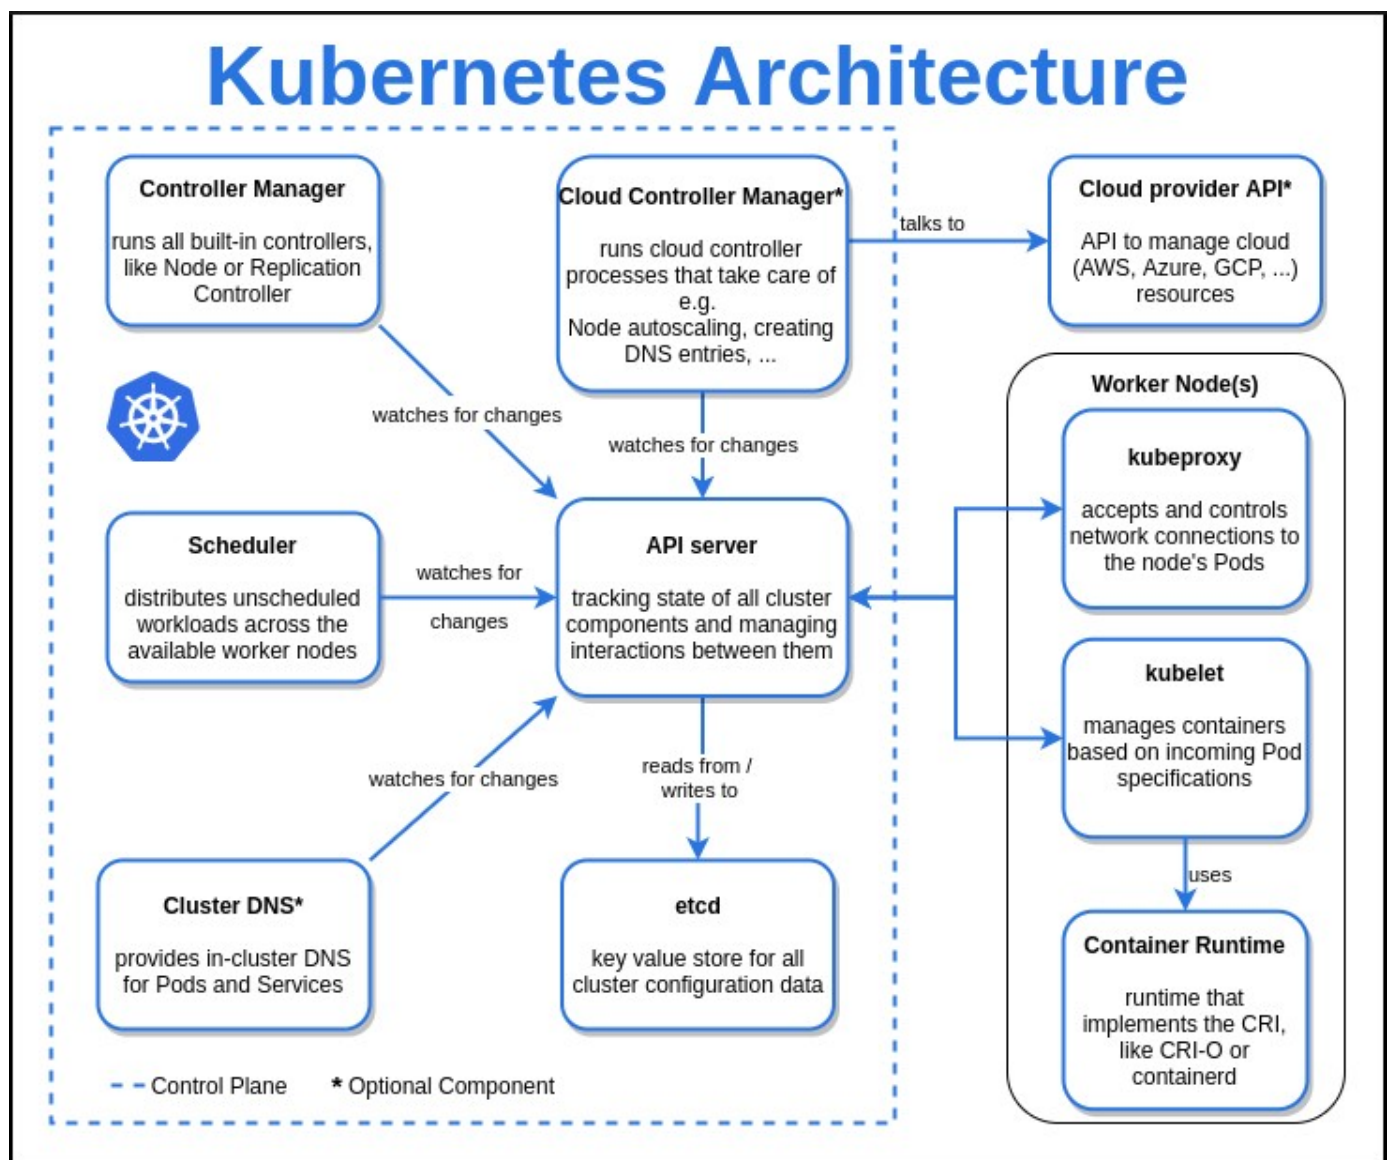
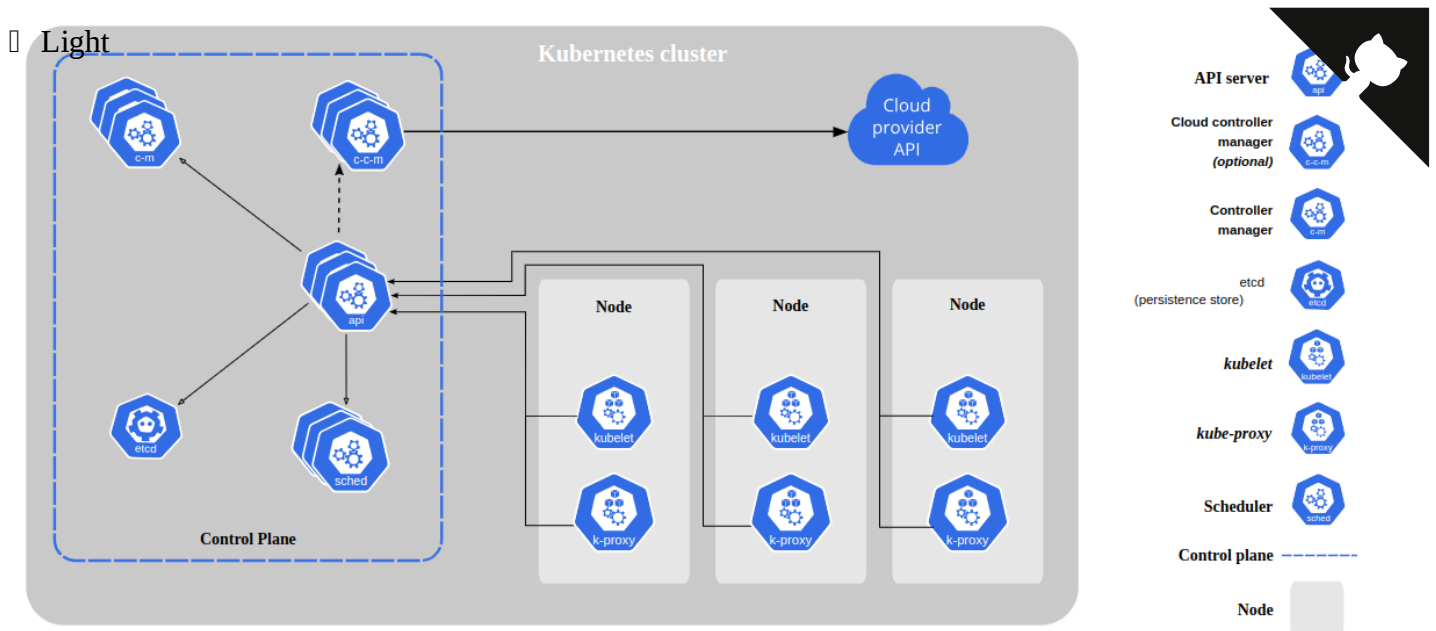




LEARNING KUBERNETES





[Explore the docs »](#)

[Main Page](#) - [Code Page](#) - [Report Bug](#) - [Request Feature](#)

Summary

► TABLE OF CONTENT

About Project

This project aims to help students or professionals to learn the main concepts of kubernetes

[\(back to top\)](#)

Getting Started

This is an example of how you may give instructions on setting up your project locally. To get a local copy up and running follow these simple example steps.

Prerequisites

This is an example of how to list things you need to use the software and how to install them.

- git
- Virtual Box and extension
- Vagrant

Installation



```
git clone https://github.com/marcoossilvestrini/learning-kubernetes.git
```

Set ssh keys in security folder

```
# generate ssh key pair for your user access hosts
ssh-keygen -q -t ecdsa -b 521 -N '' -f ~/.ssh/id_ecdsa <<<y >/dev/null 2>&1
cp ~/.ssh/id_ecdsa.pub security/

# generate ssh key pair for rancher
ssh-keygen -q -t ecdsa -b 521 -N '' -f security/kubernetes-key-ecdsa <<<y >/dev/null
```

Set Node pool

You can scale down or up number of control planes and workers in Vagrantfile array.

Example:

```
...
# Node|Control Plane Servers
PLANES = ["control-plane01", "control-plane02", "control-plane03"]
N = 2

(0..N).each do |i|
  config.vm.define PLANES[i] do |node|
  ...
```

Set network

Set network configuration for each VM in Vagrantfile.

Example:

```
...
# NETWORK
  ol9_server01.vm.network "public_network", nic_type: "virtio", mac: "080027f3066a",
    "Intel(R) I211 Gigabit Network Connection",
    "MediaTek Wi-Fi 6 MT7921 Wireless LAN"
  ]
...
```

Set VM's resources

Set resource configuration as CPU, Memory, etc for each VM in Vagrantfile.

Example:

□ Light



```
# PROVIDER
infra_server01.vm.provider "virtualbox" do |vb|
  vb.linked_clone = true
  vb.name = VM_INFRA_SERVER01
  vb.memory = 2048
  vb.cpus = 1
end
...
```

Up kubernetes cluster

```
cd learning-kubernetes/vagrant/linux
vagrant up
```

Usage

Use this repository for get learning about kubernetes exam

[\(back to top\)](#)

Roadmap

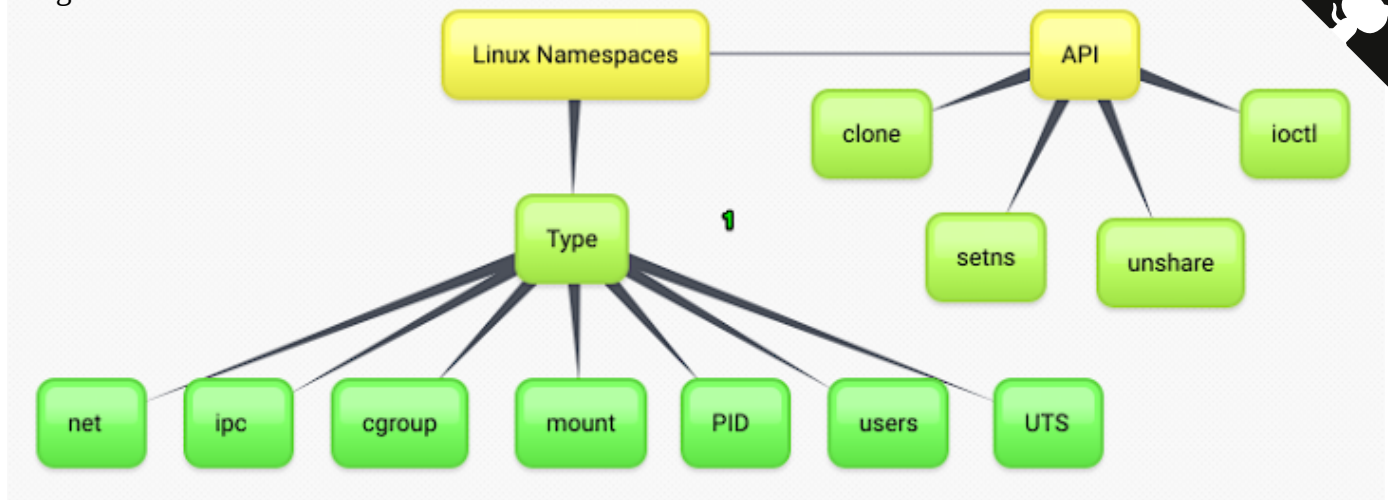
- ☒ Create repository
- ☒ Create github action for automation tasks
- ☒ Install kubernetes cluster
- ☒ Install kubectl
- ☒ Add kubernetes examples
- ☒ Add app deployment
- ☒ Create docker image with project contents
- ☒ Create github action for build docker image

[\(back to roadmap\)](#)

[\(back to top\)](#)

Linux Namespaces

kubernetes Engine work with namespaces(PID,NET,IPC,MNT,UTS) and cgroups.



pid

The Process ID. This file is a handle for the PID namespace of the process. PID namespaces isolate the process ID number space, meaning that processes in different PID namespaces can have the same PID. PID namespaces allow containers to provide functionality such as suspending/resuming the set of processes in the container and migrating the container to a new host while the processes inside the container maintain the same PIDs

net

This file is a handle for the network namespace of the process. This provides the isolation of the system resources associated with networking and isolates Network devices. The ip netns - is used to process network namespace management mount This file is a handle for the mount namespace of the process and isolates Mount points

ipc

This file is a handle for the IPC namespace of the process and isolates System Vs IPC and POSIX message queues

uts

This file is a handle for the UTS namespace of the process and isolates Hostname and NIS domain name

user

This file is a handle for the user namespace of the process



the container will have an isolated view of the cgroup hierarchy.

cgroup vs namespace

cgroup is a way to control group based traffic control filter, example

```
"cgroupsPath": "/myRuntime/myContainer",
  "resources": {
    "memory": {
      "limit": 100000,
      "reservation": 200000
    },
    "devices": [
      {
        "allow": false,
        "access": "rwm"
      }
    ]
  }
}
```

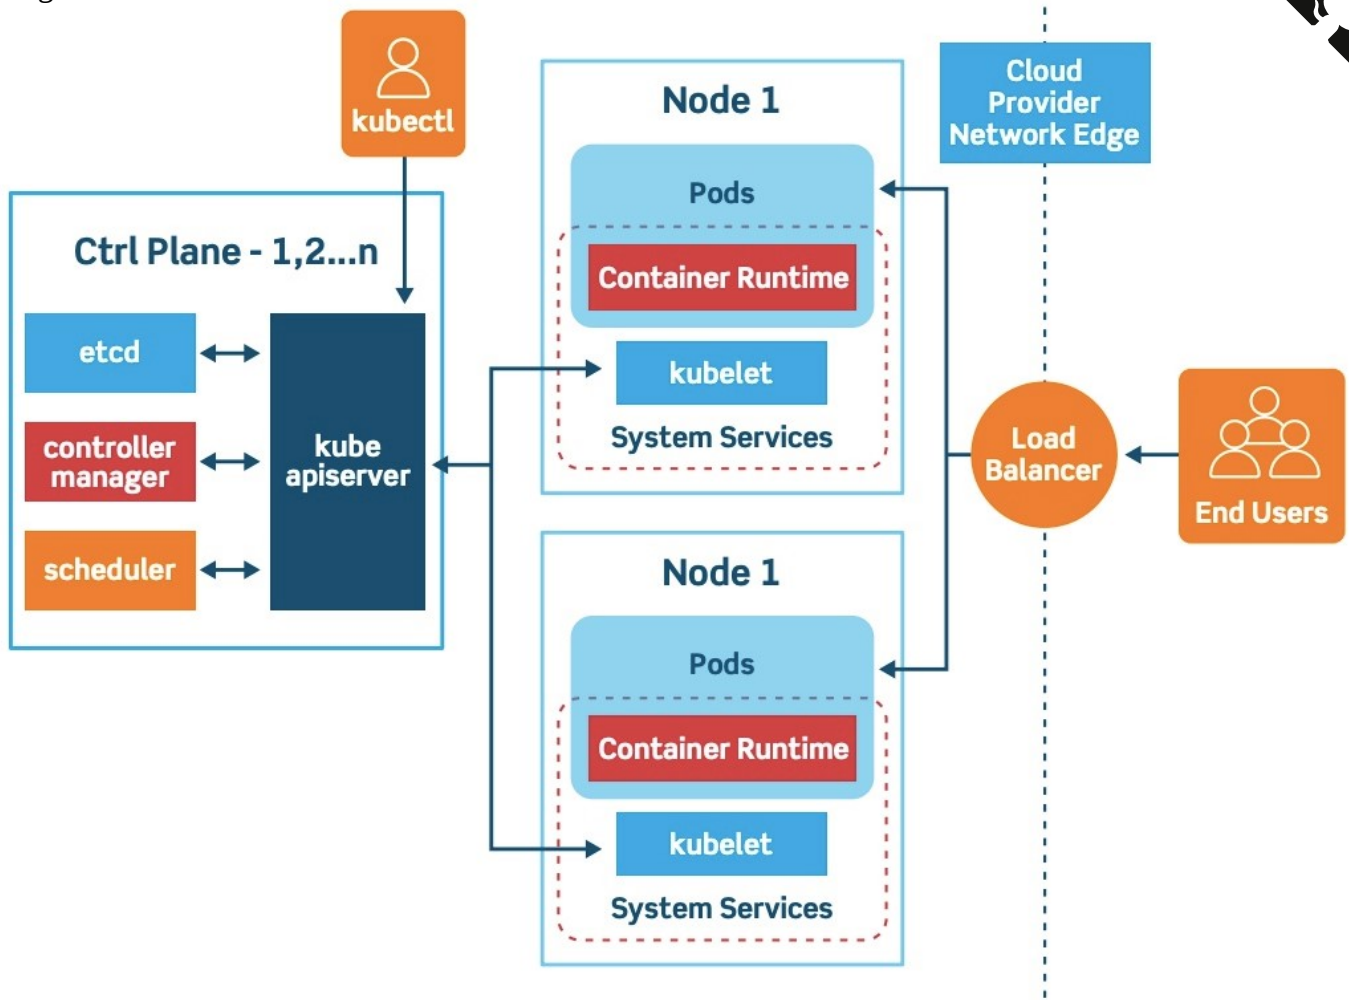
namespace: Limit/abstraction what you can see in linux proc

Font: <https://8gwifi.org/docs/linux-namespaces.jsp>


(back to linux-namespaces)

(back to top)

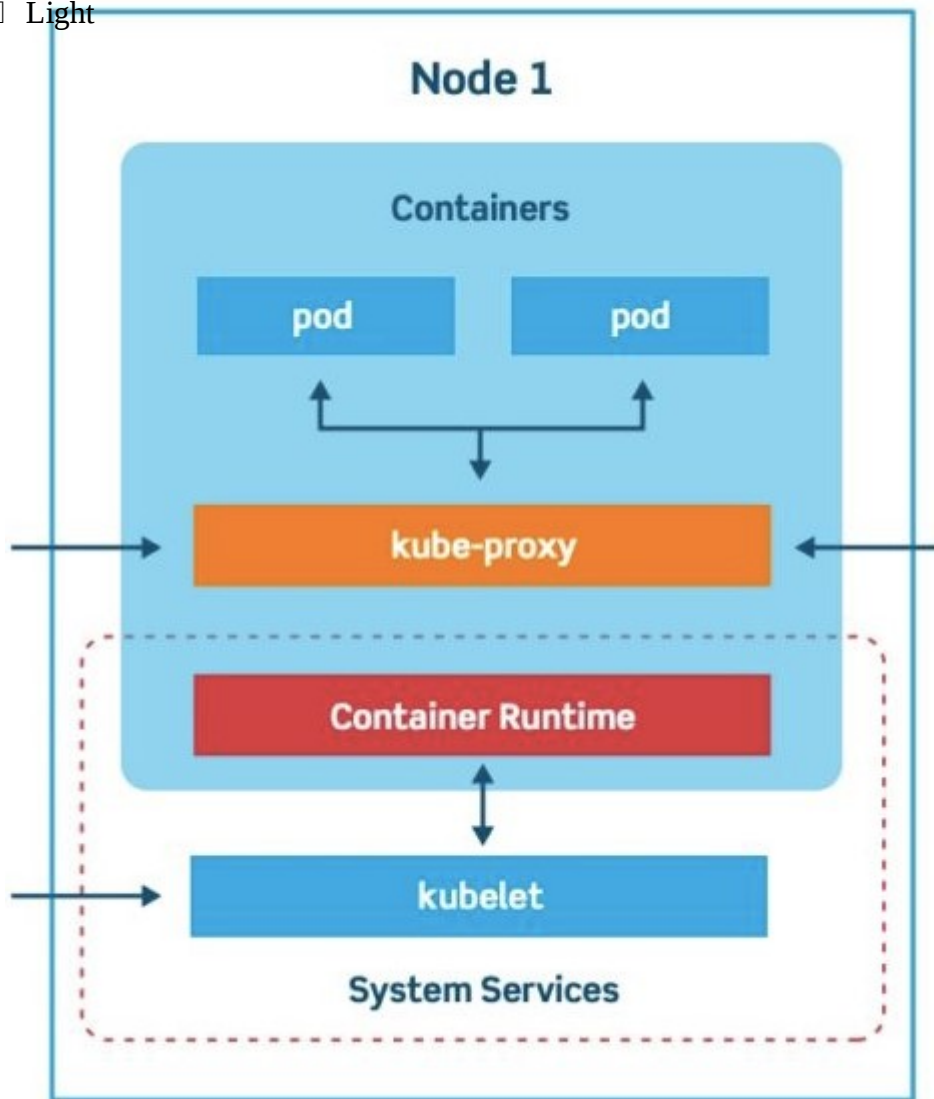
Kubernetes Architecture



Control Plane



Node



Kubernetes ports

CONTROL PLANE

Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	6443*	Kubernetes API server	All
TCP	Inbound	2379-2380	etcd server client API	kube-apiserver, etcd
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	10251	kube-scheduler	Self
TCP	Inbound	10252	kube-controller-manager	Self



Protocol	Direction	Port Range	Purpose	Used By
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	30000-32767	NodePort	Services All

Ports for Rancher Server Nodes on RKE2

Protocol	Port	Source	Description
TCP	9345	RKE2 server and agent nodes	Node registration. Port should be open on all server nodes to all other nodes in the cluster.
TCP	6443	RKE2 agent nodes	Kubernetes API
UDP	8472	RKE2 server and agent nodes	Required only for Flannel VXLAN
TCP	10250	RKE2 server and agent nodes	kubelet
TCP	2379	RKE2 server nodes	etcd client port
TCP	2380	RKE2 server nodes	etcd peer port
TCP	30000-32767	RKE2 server and agent nodes	NodePort port range. Can use TCP or UDP.
TCP	5473	Calico-node pod connecting to typha pod	Required when deploying with Calico
HTTP	80	Load balancer/proxy that does external SSL termination	Rancher UI/API when external SSL termination is used
HTTPS	443	<ul style="list-style-type: none">hosted/registered Kubernetesany source that needs to be able to use the Rancher UI or API	Rancher agent, Rancher UI/API, kubectl. Not needed if you have a load balancer doing TLS termination.

[\(back to kubernetes-architecture\)](#)

[\(back to top\)](#)

Install kubernetes

Minikube

```
# install
curl -Lo minikube https://storage.googleapis.com/minikube/releases/latest/minikube-l
```

```
Light
chmod +x ./minikube
sudo mv ./minikube /usr/local/bin/minikube

# get version
minikube version

# set hypervisor
minikube config set driver <YOUR_HYPERVISOR>

# up without hypervisor
minikube start --driver=hyperkit

# create cluster
minikube start --nodes 3 -p multinode-cluster

# get status of cluster
minikube status

# get ip address
minikube ip

# access minikube host
minikube ssh

# dashboard
minikube dashboard

# logs
minikube logs

# delete cluster
minikube delete
minikube delete --purge
```

Kind

```
# Install
curl -Lo ./kind https://kind.sigs.k8s.io/dl/v0.14.0/kind-linux-amd64
chmod +x ./kind
sudo mv ./kind /usr/local/bin/kind

# create cluster
kind create cluster
kind create cluster --name silvestrini

# get clusters
kind get clusters

# delete clusters
kind delete clusters $(kind get clusters)

## create yaml
cat << EOF > $HOME/kind-3nodes.yaml
```



```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
  - role: control-plane
  - role: worker
  - role: worker
EOF
```



```
# create cluster
kind create cluster --name kind-multinodes --config $HOME/kind-3nodes.yaml
```

[\(back to install-kubernetes\)](#)

[\(back to top\)](#)

RKE2

For create kubernetes cluster using RKE2, see scripts in folder scripts/linux/k8s/install-rke2.sh

Some commands of stack rke2

Set your PATH variable:

```
export PATH=$PATH:/opt/rke2/bin:/var/lib/rancher/rke2/bin
```

CONTAINERD - ctr Commands

```
#list containers using ctr
ctr --address /run/k3s/containerd/containerd.sock --namespace k8s.io container ls
```

CONTAINERD - crictl Commands

```
#list containers using crictl

## example 1
export CRI_CONFIG_FILE=/var/lib/rancher/rke2/agent/etc/crictl.yaml
crictl ps

## example 2
crictl --config /var/lib/rancher/rke2/agent/etc/crictl.yaml ps

## example 3
crictl --runtime-endpoint unix:///run/k3s/containerd/containerd.sock ps -a

# stats containers
crictl stats
```



```
journalctl -f -u rke2-server  
/var/lib/rancher/rke2/agent/containerd/containerd.log  
/var/lib/rancher/rke2/agent/logs/kubelet.log
```

Reference: <https://gist.github.com/superseb/3b78f47989e0dbc1295486c186e944bf>

(back to rke2)

(back to top)

Kubectl

Install

```
# install  
curl -LO https://storage.googleapis.com/kubernetes-release/release/`curl -s \  
https://storage.googleapis.com/kubernetes-release/release/stable.txt`/bin/linux/amd6  
chmod +x ./kubectl  
mv ./kubectl /usr/local/bin/kubectl  
  
# get version  
kubectl version --output=yaml --client  
  
# kubectl autocomplete  
source <(kubectl completion bash)  
  
# kubectl alias  
alias k=kubectl  
complete -F __start_kubectl k
```

Commands - Kubectl

```
##### resources #####  
  
# list all resources  
kubectl get all  
  
##### namespaces #####  
  
# get namespaces  
kubectl get namespaces  
  
# describe namespaces  
kubectl describe namespaces
```

□ Light



```
##### nodes #####
```

```
# list nodes
```

```
kubectl get nodes
```

```
kubectl get nodes -o wide
```

```
# delete node
```

```
kubectl drain <node_name> --ignore-daemonsets --delete-emptydir-data
```

```
kubectl delete node <node_name>
```

```
# get logs
```

```
kubectl logs my-nginx
```

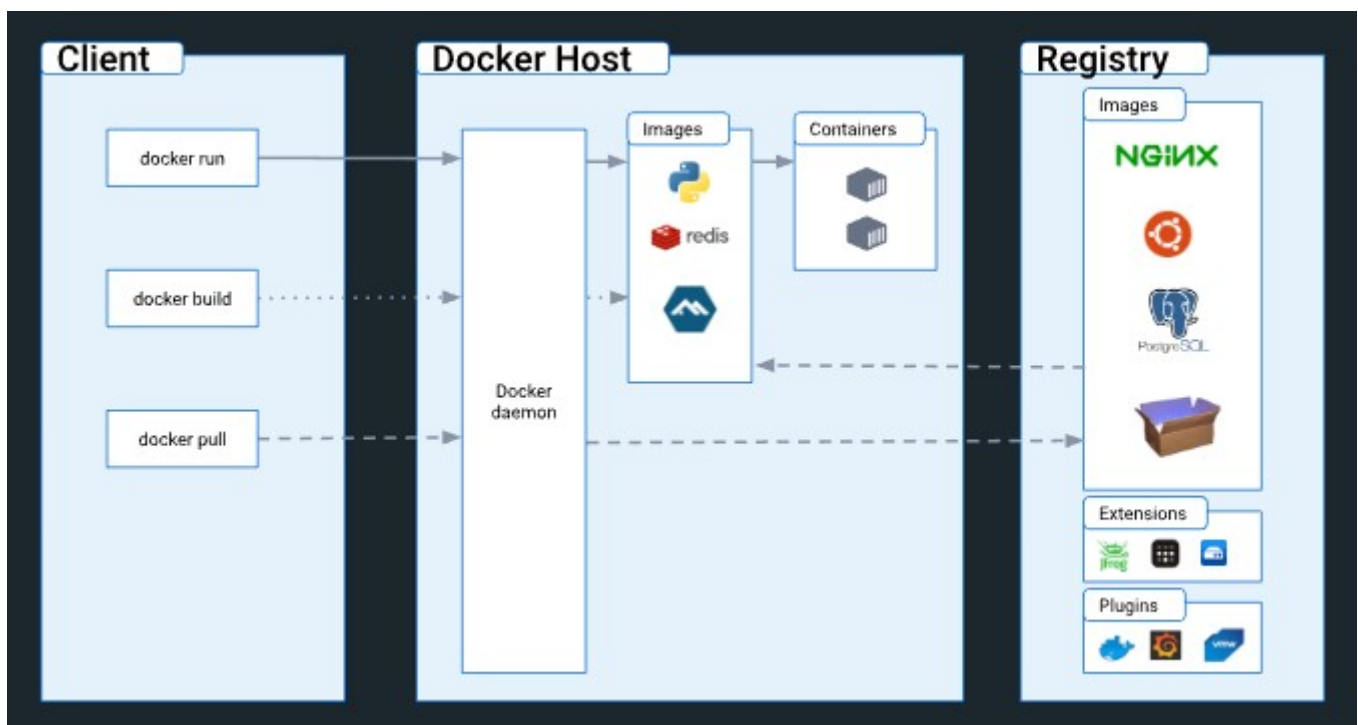
```
kubectl logs -f my-nginx
```

```
kubectl logs -n kube-system --all-containers=true etcd-control-plane01
```

(back to kubectl)

(back to top)

Containers



Commands - Containers

```
# get containers in pod
```

```
kubectl -n kube-system describe pods kube-proxy-worker01 | grep -ws -A 10 Containe
```

```
# connect in container
```

```
kubectl attach silvestrini -c infra
```

```
# connect in container
```

```
kubectl exec -it pod_name -c container_name bash
Light
kubectl exec infra ls
kubectl exec silvestrini -c infra -- ls
kubectl exec silvestrini -c infra -it sh

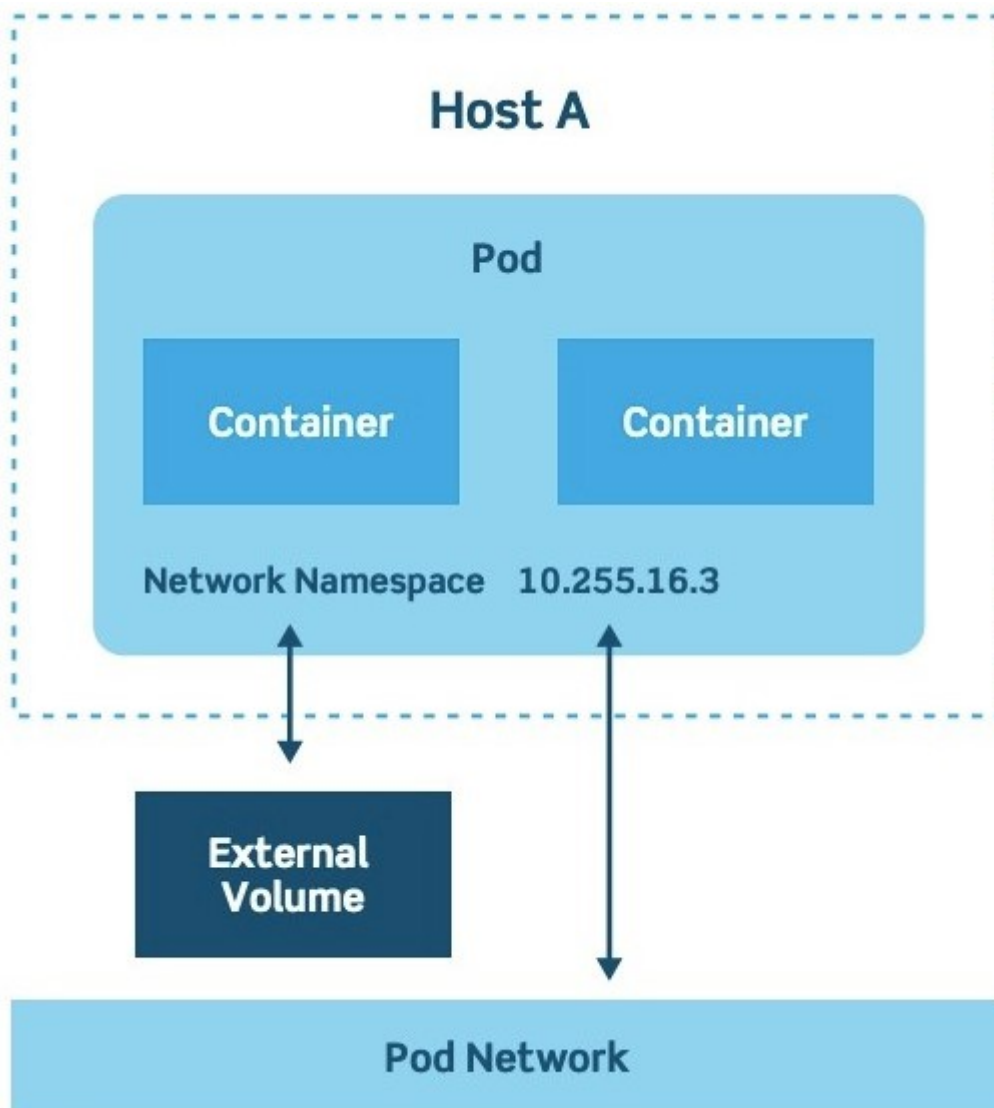
# access container in specific namespace
kubectl exec -it -n kube-system kube-proxy-worker01 -c kube-proxy -- bash
```



[\(back to containers\)](#)

[\(back to top\)](#)

Pods



A pod is the smallest execution unit in Kubernetes. A pod encapsulates one or more applications. Pods are ephemeral by nature, if a pod (or the node it executes on) fails, Kubernetes can automatically create a new replica of that pod to continue operations. Pods include one or more containers (such as Docker containers).

Pods also provide environmental dependencies, including persistent storage volumes (storage that is permanent and available to all pods in the cluster) and configuration data needed to run the container(s) within the pod.



Commands - Pods

```
# create pods without manifest
kubectl run nginx --image nginx
kubectl run -it --rm debug --image=busybox --restart=Never -- sh

# create pod with manifest
kubectl apply -f pod-template.yaml
kubectl create -f pod.yaml

# list pods
kubectl get pods

# list all pods
kubectl get pods --all-namespaces
kubectl get pods -A
kubectl get pods -A -o wide

# list pods in specific node
kubectl get pods --all-namespaces -o wide --field-selector spec.nodeName=worker01

# list pods in kube-system namespace
kubectl get pod -n kube-system
kubectl get pods -n kube-system -o=jsonpath='{range.items[*]}{"\n"}{.metadata.name}{

# list pods with specif output
kubectl get pods -n kube-system -o yaml
kubectl get pods -n kube-system -o json

# list images used in pods
kubectl get pods -o=jsonpath='{range .items[*]}{"\n"}{.metadata.name}{"\t"}{range .s

# describe details of pods
kubectl describe pod nginx
kubectl -n kube-system describe pods kube-proxy-worker01

# delete pods
kubectl delete pod nginx
kubectl delete -f pod-template.yaml


# create Service | expose pod
kubectl expose pod my-nginx
```

Understanding Pod Resources



A cluster node:

4x vCPUs 16GB RAM



1. The pod effective request
400MiB of Memory
300Mi + 100Mi
600 millicores
500m + 100m
2. Kubernetes assigns
1024 shares per core.
 $1024 \cdot 0.1 = 102 \text{ shares}$
 $1024 \cdot 0.5 = 512 \text{ shares}$

```
kind: Deployment
apiVersion: extensions/v1beta1
metadata:
  name: redis
  labels:
    name: redis-deployment
    app: example-voting-app
spec:
  replicas: 1
  selector:
    matchLabels:
      name: redis
      role: redisdb
      app: example-voting-app
  template:
    spec:
      containers:
        - name: redis
          image: redis:5.0.3-alpine
          resources:
            limits:
              memory: 600Mi
              cpu: 1
            requests:
              memory: 300Mi
              cpu: 500m
        - name: busybox
          image: busybox:1.28
          resources:
            limits:
              memory: 200Mi
              cpu: 300m
            requests:
              memory: 100Mi
              cpu: 100m
```

3. Will be killed if allocates > 600MB.
The whole Pod will fail.

4. Will be throttled if uses more than "1 Core".
1 core = 1000 millicores = 1000m =
100ms of computing time every 100 real ms
Full computing time of the node:
4 vCPUs * 100 real ms =
400ms of computing time = 4000m


5. Killed if allocates > 200MB.

6. Throttled if uses > 30ms of computing time in 100ms

(back to pods)

(back to top)

Deployment



A Deployment provides declarative updates for Pods and ReplicaSets. You describe a desired state in a Deployment, and the Deployment Controller

changes the actual state to the desired state at a controlled rate.

You can define Deployments to create new ReplicaSets, or to remove existing Deployments and adopt all their resources with new Deployments.



Commands - Deployment

```
# create manifest|template
kubectl run my-nginx --image nginx --port 80 --dry-run=client -o yaml >pod-template

# apply\update deployment
kubectl apply -f deployment.yaml

# list deployments
kubectl get deployments -A
kubectl -n kube-system get deployments.apps
kubectl get deployments -l app=nginx-deployment

# get pods management by deployment
kubectl get pods -l app=nginx-deployment

# describe deployment
kubectl describe deployment nginx-deployment

# check status of deployment
kubectl rollout status deployment nginx-deployment

# running rollback deployment
kubectl rollout undo deployment nginx-deployment
kubectl rollout undo deployment nginx-deployment --to-revision=1

# get deployment history
kubectl rollout history deployment nginx-deployment
kubectl rollout history deployment nginx-deployment --revision=2

# pause deployment(block updates)
kubectl rollout pause deployment nginx-deployment

# resume deployment(allow updates)
kubectl rollout resume deployment nginx-deployment


# restart deployment (recreate all pods in deployment)
kubectl rollout restart deployment nginx-deployment

# delete deployment
kubectl delete deployment nginx-deployment
```

[\(back to deployment\)](#)

[\(back to top\)](#)

ReplicaSet



A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time. As such, it is often used to guarantee the availability of a specified number of identical Pods.

Commands - ReplicaSet

```
# list replicaset
kubectl get replicaset -l app=nginx-deployment
```

```
# describe replicaset
kubectl describe replicaset nginx-replicaset
```

```
# create replicaset - see folder replicaset/ for examples
kubectl apply -f replicaset.yaml
```

```
# delete replicaset
kubectl delete replicas
```

[back to replicaset](#replicaset)
[back to top](#readme-top)

Daemonset

![Daemonset](images/daemonset.png)

A DaemonSet ensures that all (or some) Nodes run a copy of a Pod.\nAs nodes are added to the cluster, Pods are added to them.\nAs nodes are removed from the cluster, those Pods are garbage collected.\nDeleting a DaemonSet will clean up the Pods it created.

Commands - Daemonset

```
```sh
list daemonset
kubectl get daemonset -A

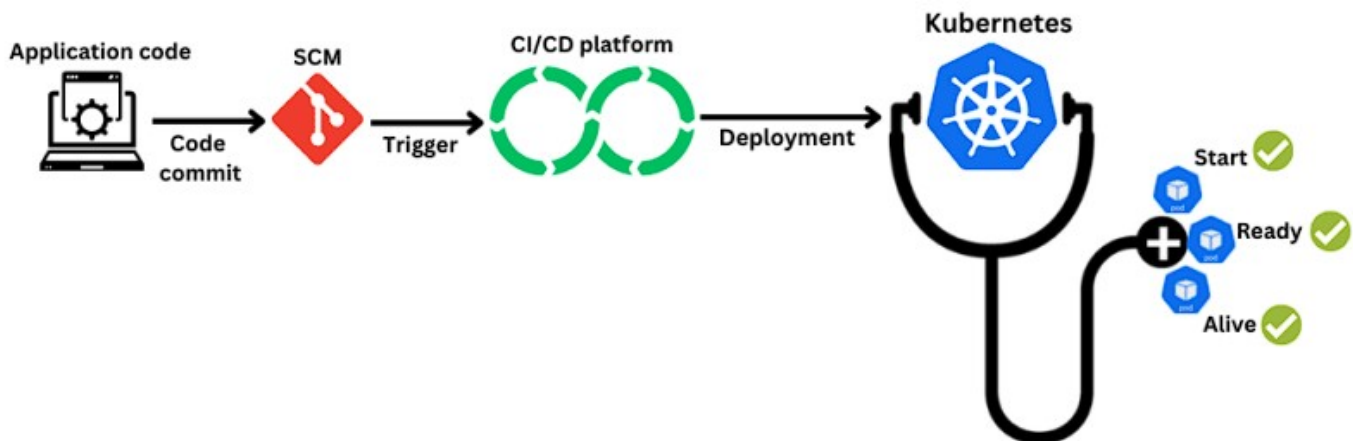
#$ describe daemonset
kubectl describe daemonset node-exporter
```

```
Light
Create daemonset - see folder daemonset/ for examples
kubectl apply -f daemonset.yaml

delete daemonset
kubectl delete daemonset node-exporter
```



## Probes



Kubernetes probes are health checks that are used to monitor the health of applications and services in a Kubernetes cluster.

Kubernetes probes are typically implemented using the Kubernetes API, which allows them to query the application or service for information.

This information can then be used to determine the application's or service's health.

Kubernetes probes can also be used to detect changes in the application or service and send a notification to the Kubernetes control plane, which can then take corrective action.

## Types of Probes

### Startup Probes

A startup probe is used to determine if a container has started successfully.

This type of probe is typically used for applications that take longer to start up, or for containers that perform initialization tasks before they become ready to receive traffic.


The startup probe is run only once, after the container has been created, and it will delay the start of the readiness and liveness probes until it succeeds.

If the startup probe fails, the container is considered to have failed to start and Kubernetes will attempt to restart the container.

### Readiness Probes

A readiness probe is used to determine if a container is ready to receive traffic.

This type of probe is used to ensure that a container is fully up and running and can accept incoming connections before it is added to the service load balancer.

Light  A readiness probe can be used to check the availability of an application's dependencies or to perform any other check that indicates the container is ready to serve traffic. If the readiness probe fails, the container is removed from the service load balancer until the probe succeeds again.

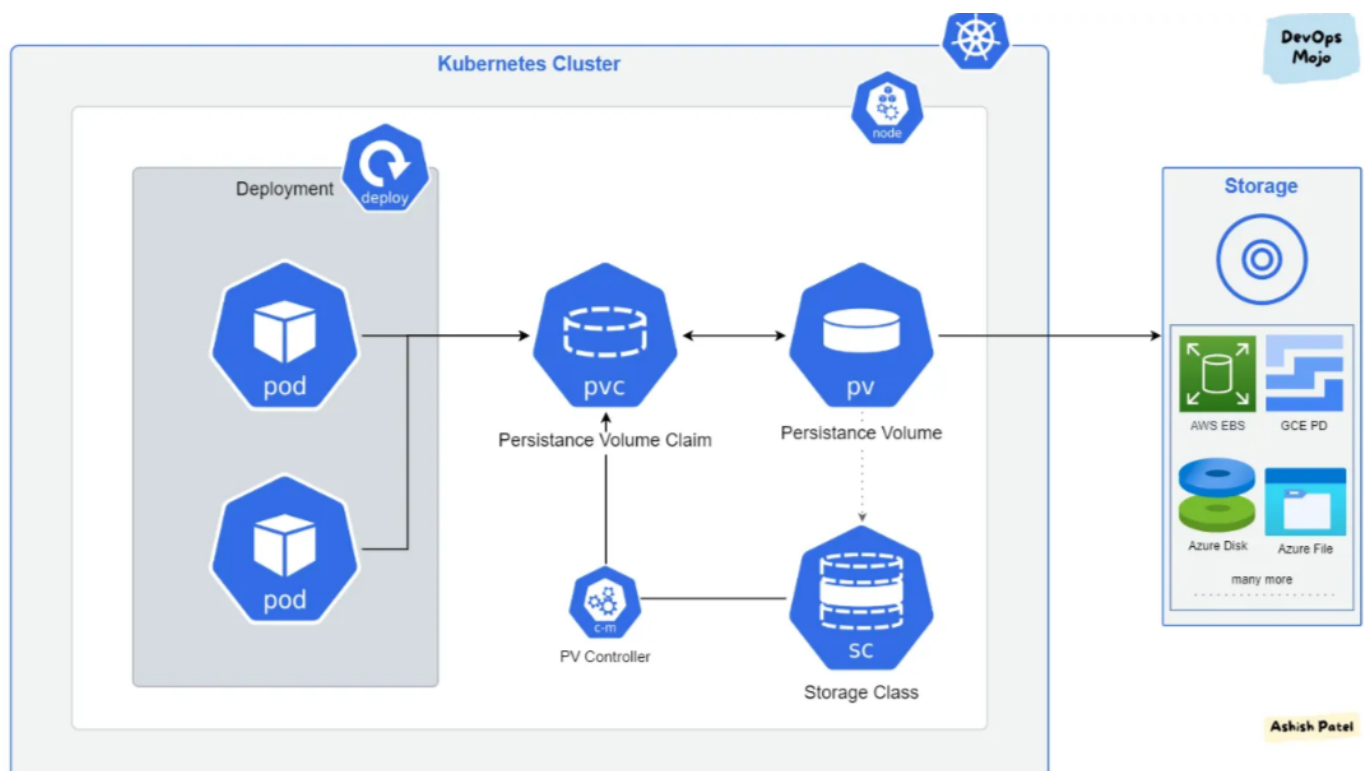
## Liveness Probes

A liveness probe is used to determine if a container is still running and functioning properly. This type of probe is used to detect and recover from container crashes or hang-ups. A liveness probe can be used to check the responsiveness of an application or to perform any other check that indicates the container is still alive and healthy. If the liveness probe fails, Kubernetes will attempt to restart the container to restore its functionality.

[\(back to probes\)](#)

[\(back to top\)](#)

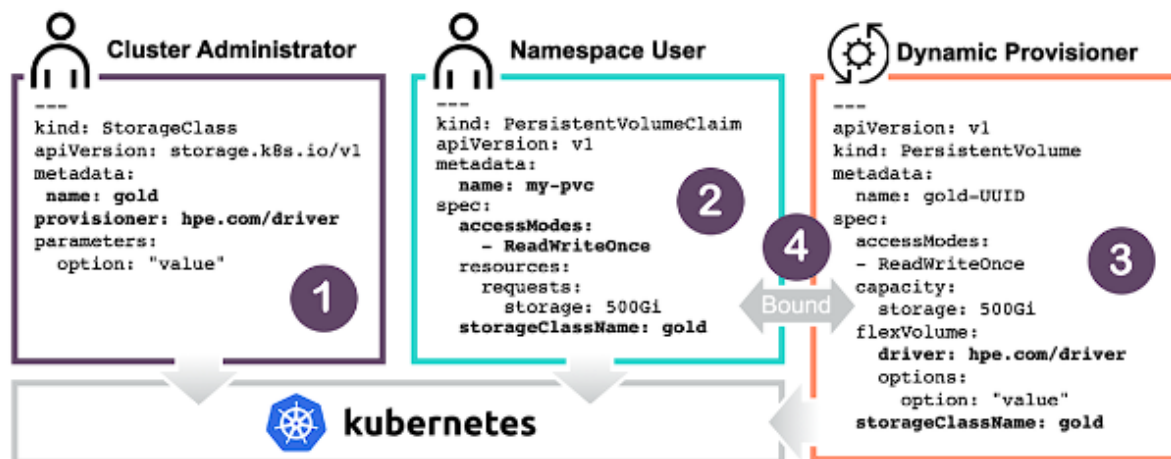
## Volumes



## Storage Class



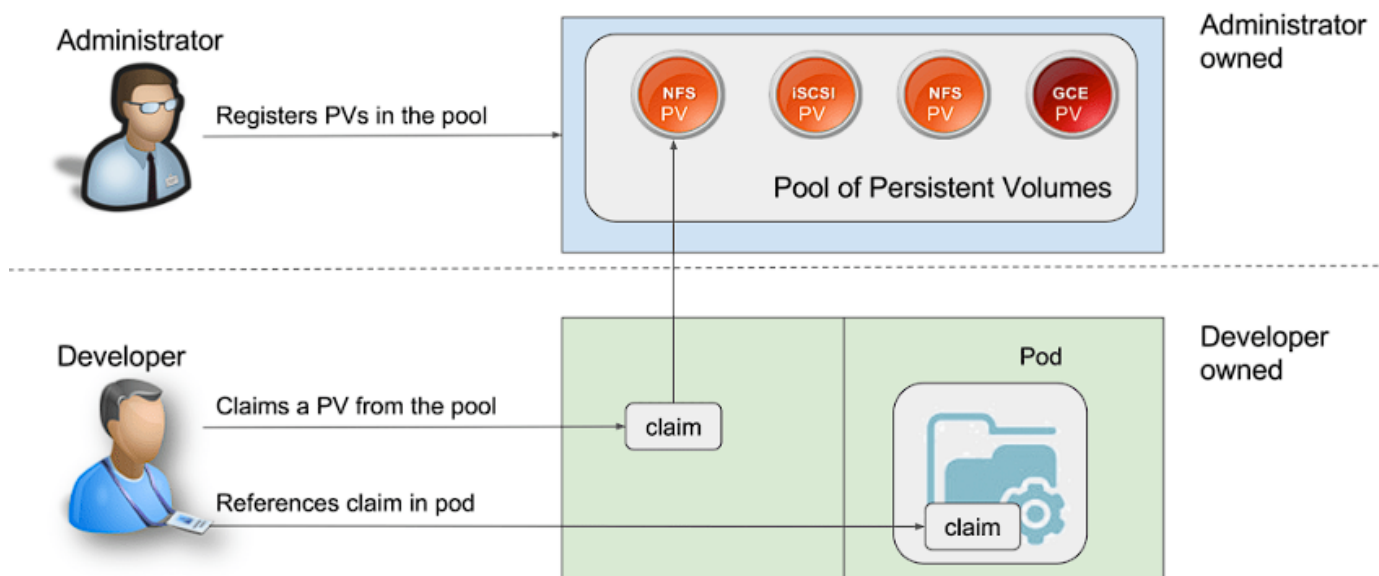
# Dynamic Provisioning for Kubernetes



Some Storage Class providers

- `kubernetes.io/aws-ebs`: AWS Elastic Block Store (EBS)
- `kubernetes.io/azure-disk`: Azure Disk
- `kubernetes.io/gce-pd`: Google Compute Engine (GCE) Persistent Disk
- `kubernetes.io/cinder`: OpenStack Cinder
- `kubernetes.io/vsphere-volume`: vSphere
- `kubernetes.io/no-provisioner`: Volumes locaux
- `kubernetes.io/host-path`: Volumes locaux

## PV - Persistent volume



Some Types of PV

## □ Light

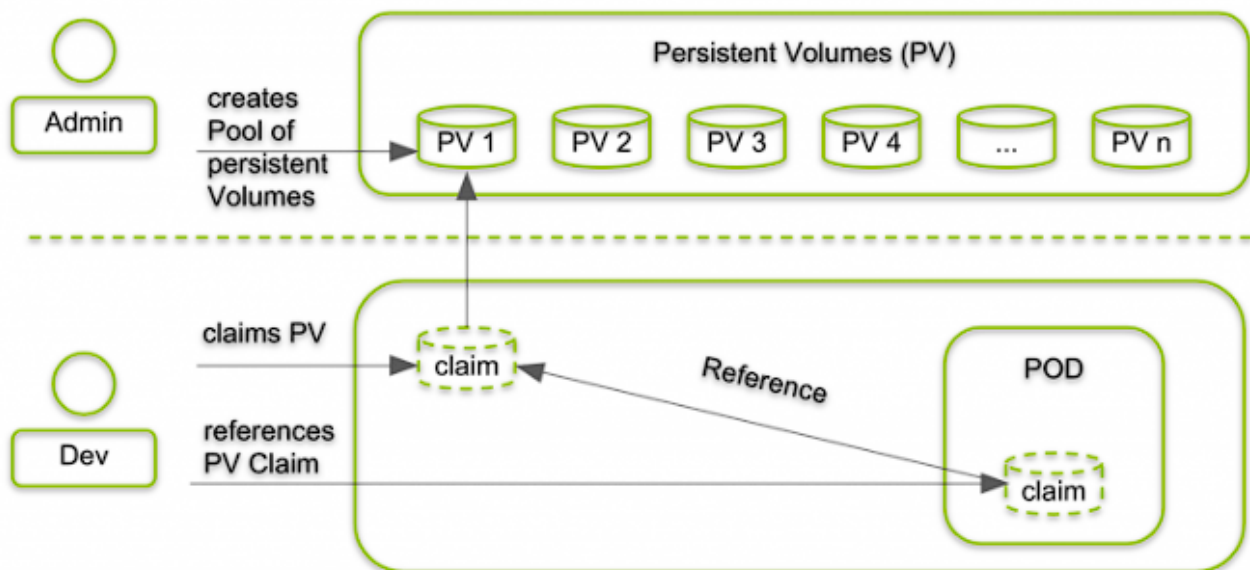


- Local
  - HostPath
- Network
  - NFS
  - iSCSI
  - Ceph RBD(RAidos Block Device)
  - GlusterFS
  - Cloud Providers(EBS,Google Cloud Persistent Disk,Azure Disk Storage)

### Type of Storage(hostPath)

- hostPath
- nfs
- iscsi
- csi
- local
- fc

### PVC - Persistent Volume Claim



### Commands - Volumes

```
list storage classes
kubectl get storageclass

describe storage class
kubectl describe storageclass silvestrini
```

Light

```
get storage class provisioners
kubectl get storageclasses.storage.k8s.io -o=jsonpath='{range.items[*]}.{.provisioner}'

list pv in cluster
kubectl get pv -A

describe pv
kubectl describe pv my-pv

list pvc
kubectl get pvc -o wide

delete pvc
kubectl delete pvc my-pvc

describe pvc
kubectl describe pvc my-pvc


get events
kubectl get events my-pvc.1772cda2d4c7069b
```

[\(back to volumes\)](#)

[\(back to top\)](#)

## StatefulSet





StatefulSet is the workload API object used to manage stateful applications.

Manages the deployment and scaling of a set of Pods, and provides guarantees about the ordering and uniqueness of these Pods.

Like a Deployment, a StatefulSet manages Pods that are based on an identical container spec. Unlike a Deployment, a StatefulSet maintains a sticky identity for each of its Pods.\

These pods are created from the same spec, but are not interchangeable: each has a persistent identifier that it maintains across any rescheduling.

If you want to use storage volumes to provide persistence for your workload, you can use a StatefulSet as part of the solution. Although individual Pods in

□ Light  
a StatefulSet are susceptible to failure, the persistent Pod identifiers make it easier to match existing volumes to the new Pods that replace any that have failed.



## DNS for pods in StatefulSet

```
<pod-name>.<service-name>.<namespace>.svc.cluster.local
Example: nginx-0.nginx.default.svc.cluster.local
```

## Commands - StatefulSet

```
list statefulsets
kubectl get statefulsets

describe statefulsets
kubectl describe statefulsets.apps nginx

delete statefulset
kubectl delete statefulset nginx

test network service
create container for test
kubectl run -it --rm --image=busybox --restart=Never -- sh


test nslookup
nslookup nginx-0.nginx.default.svc.cluster.local

test web page
wget -O- http://nginx-0.nginx.default.svc.cluster.local
```

[\(back to statefulset\)](#)

[\(back to top\)](#)

## Services



A Kubernetes service is a logical abstraction for a deployed group of pods in a cluster (which all perform the same function).

Since pods are ephemeral, a service enables a group of pods, which provide specific functions (web services, image processing, etc.) to be assigned a name and unique IP address (clusterIP).

## Service Types in Kubernetes

### ClusterIP Services

ClusterIP is the default service type in Kubernetes, and it provides internal connectivity between different components of our application. Kubernetes assigns a virtual IP address to a ClusterIP service that can solely be accessed from within the cluster during its creation. This IP address is stable and doesn't change even if the pods behind the service are rescheduled or replaced.

ClusterIP services are an excellent choice for internal communication between different components of our application that don't need to be exposed to the outside world. For example, if we have a microservice that processes data and sends it to another microservice for further processing, we can use a ClusterIP service to connect them.

To create a ClusterIP service in Kubernetes, we need to define it in a YAML file and apply it to the cluster. Here's an example of a simple ClusterIP service definition:

```
apiVersion: v1
kind: Service
metadata:
 name: backend
spec:
 selector:
 app: backend
```

```
ports:
- name: http
 port: 80
 targetPort: 8080
```



In this example, we define a service named backend with a selector that targets pods labeled with app: backend. The service exposes port 80, which is the port used by clients to access the service, and forwards the traffic to the pods' port 8080, which is where the backend application is running.

## NodePort Services

NodePort services extend the functionality of ClusterIP services by enabling external connectivity to our application. When we create a NodePort service on any node within the cluster that meets the defined criteria, Kubernetes opens up a designated port that forwards traffic to the corresponding ClusterIP service running on the node.

These services are ideal for applications that need to be accessible from outside the cluster, such as web applications or APIs. With NodePort services, we can access our application using the node's IP address and the port number assigned to the service.

Let's look at an example of a simple NodePort service definition:

```
apiVersion: v1
kind: Service
metadata:
 name: frontend
spec:
 selector:
 app: frontend
 type: NodePort
 ports:
 - name: http
 port: 80
 targetPort: 8080
```

We define a service named frontend that targets pods labeled with app: frontend by setting a selector. The service exposes port 80 and forwards the traffic to the pods' port 8080. We set the service type to NodePort, and Kubernetes exposes the service on a specific port on a qualifying node within the cluster.

When we create a NodePort service, Kubernetes assigns a port number from a predefined range of 30000-32767. Additionally, we can specify a custom port number by adding the nodePort field to the service definition:

```
apiVersion: v1
```

```
kind: Service
Light metadata:
 name: frontend
spec:
 selector:
 app: frontend
 type: NodePort
 ports:
 - name: http
 port: 80
 targetPort: 8080
 nodePort: 30080
```



The nodePort field is specified as 30080, which tells Kubernetes to expose the service on port 30080 on every node in the cluster.

## LoadBalancer Services

LoadBalancer services connect our applications externally, and production environments use them where high availability and scalability are critical. When we create a LoadBalancer service, Kubernetes provisions a load balancer in our cloud environment and forwards the traffic to the nodes running the service.

LoadBalancer services are ideal for applications that need to handle high traffic volumes, such as web applications or APIs. With LoadBalancer services, we can access our application using a single IP address assigned to the load balancer.

Here's an example of a simple LoadBalancer service definition:

```
apiVersion: v1
kind: Service
metadata:
 name: web
spec:
 selector:
 app: web
 type: LoadBalancer
 ports:
 - name: http
 port: 80
 targetPort: 8080
```

We set the service type to LoadBalancer to instruct Kubernetes to provision a load balancer. Here, we define a service named web and specify a selector that targets pods labeled with app: web. Additionally, we expose port 80 and forward traffic to the pods' port 8080.

After creating the LoadBalancer service, Kubernetes provisions a load balancer in the cloud environment with a public IP address. We can use this IP address to access our application from outside the cluster.



Service Type	Use Case	Accessibility	Resource Allocation
<i>ClusterIP</i>	Internal communication between application components	Within the cluster only	Minimal resources needed
<i>NodePort</i>	External accessibility for web applications or APIs	Accessible from outside the cluster via a high-numbered port on the node	Additional resources needed
<i>LoadBalancer</i>	Production environments with high traffic volumes	Accessible from outside the cluster via a load balancer	Significant resources needed
Cloud Provider's Load Balancer	Using a cloud provider for Kubernetes	Accessible from outside the cluster via the cloud provider's load balancer	May result in cost savings and better performance

Reference: <https://www.baeldung.com/ops/kubernetes-service-types>

## Commands - Services

```
list services
kubectl get services
kubectl get svc -o wide

list services in system namespace
kubectl get svc -n kube-system

details of services
kubectl describe svc nginx
kubectl describe svc -n kube-system
```

```
list endpoints
Light
kubectl get endpoints

create ClusterIP service
kubectl expose deployment app-silvestrini --port=80 --target-port=8080

create NodePort service
kubectl expose deployment app-silvestrini --type=NodePort --port=80 --target-port=80

create loadbalance service
kubectl expose deployment app-silvestrini --type=LoadBalancer --port=80 --target-port=80


create external name service
kubectl create service externalname app-silvestrini --external-name my-db.skynet.com

delete service
kubectl delete service nginx
```

[\(back to services\)](#)

[\(back to top\)](#)

## Secrets



A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key. Such information might otherwise be put in a Pod specification or in a container image.

Using a Secret means that you don't need to include confidential data in your application code.



Because Secrets can be created independently of the Pods that use them, there is less risk of the Secret (and its data) being exposed during the workflow of creating, viewing, and editing Pods. Kubernetes, and applications that run in your cluster, can also take additional precautions with Secrets, such as avoiding writing secret data to nonvolatile storage.

Secrets are similar to ConfigMaps but are specifically intended to hold confidential data.

## Types of Secrets

- Opaque Secrets - these are the simplest and most common secrets. They store arbitrary data such as API keys, passwords, and tokens. Opaque Secrets are encoded in base64 when stored in Kubernetes, but they are not encrypted. They can be used to store sensitive data but are not secure enough for highly sensitive information like database passwords.
- `kubernetes.io/service-account-token` - used to store service account access tokens. These tokens are used to authenticate Pods with the Kubernetes API. They are automatically mounted in Pods that use service accounts.
- `kubernetes.io/dockercfg` and `kubernetes.io/dockerconfigjson` - used to store Docker registry credentials. They are used to authenticate Pods with a Docker registry. They are mounted in Pods that use private container images.
- `kubernetes.io/tls`, `kubernetes.io/ssh-auth`, and `kubernetes.io/basic-auth` - used to store TLS certificates, SSH keys, and basic authentication credentials, respectively. They are used to authenticate Pods with other services.
- `bootstrap.kubernetes.io/token` - used to store cluster bootstrapping tokens. They are used to authenticate nodes with the Kubernetes control plane.

## Commands - Secrets

```
get secrets
kubectl get secrets -A

describe secrets
kubectl describe secret -n cert-manager cert-manager-webhook-ca

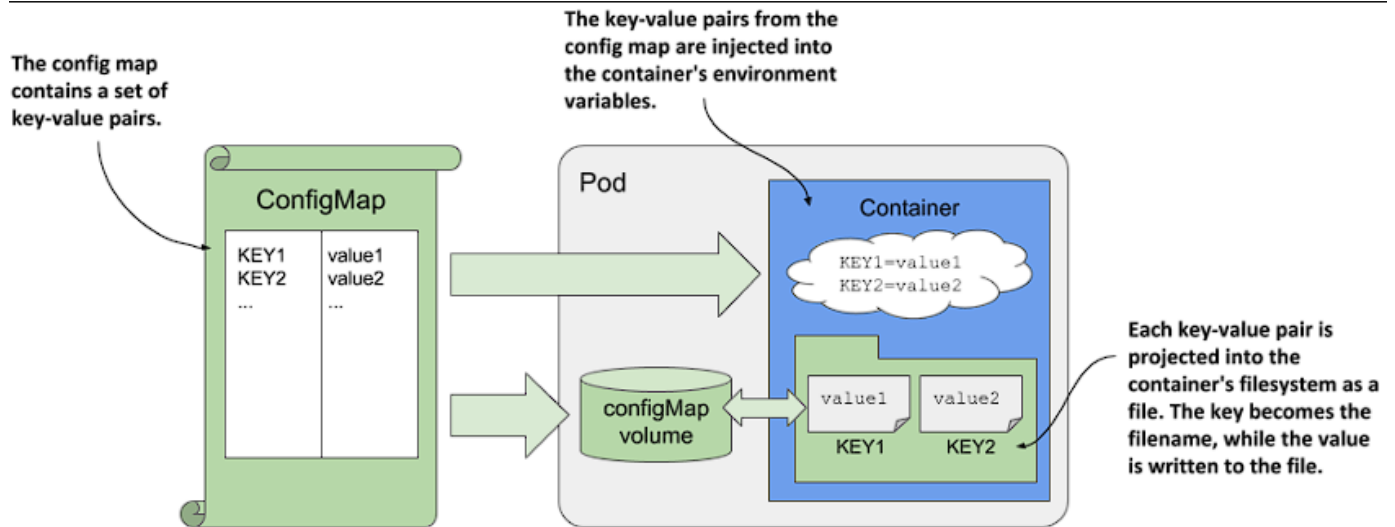
create opaque secret
kubectl create secret generic silvestrini-secret --from-literal=username=silvestrini

create tls secret
generate certs
openssl req -x509 -nodes -days 365 -newkey rsa:2048 -keyout private-key.key -out c
kubectl create secret tls my-service-web-tls-secret --cert=cert.crt --key=private-ke
```

[\(back to secrets\)](#)



# ConfigMaps



A ConfigMap is an API object used to store non-confidential data in key-value pairs. Pods can consume ConfigMaps as environment variables, command-line arguments, or as configuration files in a volume.

A ConfigMap allows you to decouple environment-specific configuration from your container images, so that your applications are easily portable.

## Commands - ConfigMaps

```
create configmap
kubectl create configmap nginx-config --from-file=kubernetes/configmaps/nginx.conf

get configmaps
kubectl get configmaps

describe configmaps
kubectl describe configmaps nginx-config
```

[\(back to configmap\)](#)

[\(back to top\)](#)

## Contributing

Contributions are what make the open source community such an amazing place to learn, inspire, and create. Any contributions you make are **greatly appreciated**.

If you have a suggestion that would make this better, please fork the repo and create a pull request. You can also simply open an issue with the tag "enhancement". Don't forget to give the project a star! Thanks again!



1. Fork the Project
2. Create your Feature Branch ( `git checkout -b feature/AmazingFeature` )
3. Commit your Changes ( `git commit -m 'Add some AmazingFeature'` )
4. Push to the Branch ( `git push origin feature/AmazingFeature` )
5. Open a Pull Request

## License

- This project is licensed under the MIT License \* see the LICENSE.md file for details

## Contact

Marcos Silvestrini - [marcos.silvestrini@gmail.com](mailto:marcos.silvestrini@gmail.com)



Project Link: <https://github.com/marcoossilvestrini/learning-kubernetes>

(back to top)

## Acknowledgments

- CNCF - Cloud Native Computing Foundation
- OCI - Open Container Initiative
- Borg
- kubernetes Website
- Kubernetes Architecture
- Github
- Issues
- CKA Certification
- CKAD Certification
- CKS Certification
- Kind
- Minikube
- k0s
- k3s
- RKE2



## Up RKE2 Cluster HA

- Kubernetes Workloads
- Rancher Local Path Provisioner
- Kubernetes Dashboard
- Openlens Cluster Managment
- Book Linuxips
- Kubernetes Services



[\(back to top\)](#)